


Center Circle

Project Overview & Guidelines



Change Log

Date	Updates
06/15	Updates sections: Per-Task Loop (New UI), HOW TO CREATE YOUR COMPARISON CODE , The Three Golden Rules of Multi-Turn Chats
06/09	Link to Outlier Playground added; How to Generate & Paste Shareable Links , section improved; “5. Running the Conversation — Turns & Stop Condition” improved
06/08	Section “4. Per-Task Loop” updated.

1. Project Introduction

Center Circle is a product-experience evaluation comparing how leading AI platforms handle the same task over a full conversation. Each task gives you a **preset prompt**, a set of **Considerations for Next Steps**, and a **final goal to reach**. You run that prompt against **two** assigned models, carry each conversation as far as it needs to go **(up to 10 turns)** to reach a fully satisfying answer, score each model on 11 dimensions, write feedback for each, and then pick which model you preferred.

⚠ CRITICAL REQUIREMENT: To perform evaluations on this project, you must have accounts on the external AI platforms assigned to your queue (e.g., Claude, ChatGPT, Gemini, etc.). While a paid subscription is NOT required, please keep in mind that these tasks can be quite detailed and may require up to 10 conversation turns per model. As a result, it is possible to reach the limits of a free-tier account. If this happens, please pause your work and return to the tasking once your usage limit has been reset.

Your scores and written feedback are the data this project depends on.

⚠ WARNING - NO LLM-WRITTEN FEEDBACK: Large Language Models are strictly prohibited for writing your feedback or justifications. All written evaluations must be your original analysis. Using AI to generate these responses compromises data integrity and will result in immediate removal from the project.

🌐 LANGUAGE NOTE: The prompt is delivered ready-to-paste in your selected language - you do **not** translate it yourself. However, **all feedback and justifications must be written in English.**

2. The Three Task Elements

Every task shows you three things up front. Read all three **before** you run anything - together they tell you what a successful conversation looks like and when to stop.

Element	What it is	How you use it
Prompt	The exact user message, ready to paste in your language.	Copy it into the tool to generate the comparison code, then run it against each model. Do not rewrite it.
Considerations for Next Steps/Recommended Conversation Flow	A short list of plausible prompts to use to continue the conversation.	Your follow-up turns must come from this list . Once you've used those prompts, you can continue the conversation with your own prompts if necessary.
Final goal to reach	A plain statement of what a <i>fully satisfying answer</i> means for this task.	This is your stop condition . Keep going until the conversation satisfies it; then stop.

3. Setup - Install the Outlier Playground Extension (*one-time*)

This Chrome extension is your primary interface for session recording and response comparison. You cannot access tasks until installation is verified.

- **Install:** Add the Outlier Playground extension from the [Chrome Web Store](#).
- **Pin it:** Click the puzzle-piece icon in your Chrome toolbar and pin Outlier Playground for easy access.
- **Verify:** Confirm the icon is visible in your toolbar before moving on.

4. Per-Task Loop (New UI) [WORKFLOW VIDEO](#)

Phase 1: Setup and Code Generation

- **Extension:** Ensure you have the Outlier Playground extension installed from the [Chrome Web Store](#).
- **Launch the Extension:** After claiming a task, click the puzzle piece icon located in the top right corner of your Chrome browser, and then click on **Outlier Playground**.
- **Copy Your Given Prompt:** Copy your preset task prompt from the left dashboard under the instructions.
- **Enter Comparison Code:** Scroll down to the showdown tool on the main dashboard. Paste your prompt text into the input field and click **Run Tool** to generate your code. Copy this code, paste it into the "Enter your comparison code" field in the extension panel on the right, and click **Continue**.
- **Accept Permissions:** A disclaimer will pop up letting you know that a recording of the session will be saved (recorded in-browser only). Press **Allow & Continue**.

Phase 2: First Model Evaluation

- **Navigate to the Model:** Click the provided link in the extension interface to launch the first product you will be testing in a new tab.
- **Generate the Response:** The extension plugin will attempt to auto-fill your starting prompt and run the first turn automatically in the background. Please allow the model to finish the process automatically. If that doesn't happen (in a rare case), manually copy and paste your prompt to generate the response. Wait patiently for the model to finish its response.
- **Assess the Response:** Read the response and take notes on its formatting, styling, layout, and how well it targets the project goal.
- **Log the Next Step:** Scroll down to the bottom of the main dashboard. Your first follow-up turn (Turn 2) must align with the provided guidelines:
 - **Core Flow Requirement:** You must use these recommended follow-up questions (under the 'Recommended Conversation Flow' list) in order.

Recommended Conversation Flow

Use this list of follow-up prompts that are taken from the "Consideration of Next Steps" after your initial response to continue the model conversation.

1. Can you add an open book and a stack of novels next to the cup of tea?
2. Can you add a sleeping cat or a small dog curled up on the blanket in the reading nook?
3. Can you make the aspect ratio wider so the image works as a desktop wallpaper?

- Use Prompt #1 from the list for Turn 2, Prompt #2 from the list for Turn 3, and Prompt #3 from the list for Turn 4. **Try to continue the conversation for as many turns as possible (up to 10 turns). However, if the model reaches the goal earlier, you may end the conversation at that point.**
- **In the chat:** Copy the exact first prompt from the list in YOUR local language (for non-English contributors) and paste it verbatim directly into your conversation with the model.
- **On the dashboard:** Copy and paste that exact same verbatim prompt in **English** into the "Selected next step" field on your dashboard.
- **Extending the Conversation:** Once you have used all of the prompts from the **Recommended Conversation Flow**, you may continue with additional follow-up questions if needed. To do this, you can either use the remaining 2 follow-up prompts taken from the "Considerations for Next Steps" list or come up with your own natural questions.
- **Assess the rest of the conversation:** Read and assess the rest of the responses carefully.
- **Document Your Experience:** Write your evaluation directly in the extension panel on the right (under the "Notes on Model's response" field). Explain your thoughts on the product experience, ensuring you outline specific strengths and weaknesses.
- **Transfer to Outlier & Score:**
 - Copy your written justification from the extension panel and paste it into the **Model Feedback** field on the Outlier dashboard. **DO NOT click the "Done with conversation" button in the extension panel until you have successfully copied your written feedback. Once you click "Done with conversation" you will not be able to reopen that extension screen or access that text again. Always copy your written justification from the extension panel and paste it into the Model Feedback field on the main Outlier dashboard before completing the session.**
 - Create and copy the shareable link to the conversation from the model.
 - Paste this link into the **Model & Published Conversation Link** field on the Outlier main dashboard.
 - Indicate the total **Number of turns** completed for this model.
 - Score the response across the 11 dimensions.
- **Complete the Session:** Return to the extension panel and click **Done with conversation**. This saves your progress and automatically redirects you to the next model steps.

Phase 3: Second Model Evaluation & Final Comparison

- **Evaluate the Second Model:** Launch the second model link from the extension interface. Follow the exact same flow (**matching steps from Phase 2**) to maintain perfectly parallel evaluations. Note that if a model fails to address key constraints (like budget or location), you may need to add extra follow-up turns to reach the core goals.

- **Select Your Preference:** After completing the second evaluation, logging your feedback, and clicking **Done with conversation**, the extension panel will ask **Which model do you prefer?**
- **Select Useful UI Features:** Check off the specific UI elements you found helpful during the evaluations (such as Interactive maps, Tables, etc.). Ensure these selections match your written text precisely so that the automated back-end linters do not flag your submission.
- **Justify Your Choice:** Under the **Final Justification** section, paste your overall comparison paragraph. This is the only place where you should actively compare the two models against each other.
- **Submit:** Click **Submit comparison** within the extension panel, and then click the blue **Submit task** button at the top right corner of your main dashboard to complete your work.


**REFER TO THE FOLLOWING SCREENSHOTS TO
LEARN HOW TO CREATE YOUR COMPARISON CODE
AND WHERE TO WRITE YOUR FIRST (PRESET)
PROMPT!**

For the very 'first' prompt, please follow the below step-by-step process.

DO NOT rewrite the first prompt, translate it yourself, or inject unauthorized instructions (such as hidden scoring rules) into the Outlier Playground Extension.


1 Step 1: Open the task, and check the provided Preset Prompt

(This is set in your local language → e.g., If you are working on a ko_KR task, this preset prompt is in Korean / if you are working on the en_US task, it is in English, etc.)




- → Here, since it is an en_US task, the preset prompt is in English.


2 Step 2: Scroll down and 'Click create-show-code from the drop and press Add Tool.



3 Step 3: Paste your preset prompt into the 'Create-showdown-code Input' box → Click 'Run Tool' to generate the code → A code will be generated. → Copy the code into the Outlier Playground Extension



4 Step 4: Open the Outlier Playground Extension, and paste the generated code into the right-hand sidebar.



5 Step 5: After entering into the code, click 'Continue' → click 'Allow & Continue' (here, you have to allow the recording permission) → Once you've finished these steps, a model name box will be shown → Click it → The tool will guide you to open the first model tab → The model (whichever it is) will then 'automatically' fill in the preset prompt and a response.


Wait until the model auto-fill the preset prompt & response. DO NOT try to type in the prompt manually.

Repeat this process for the next model.

How to Generate & Paste Shareable Links

 **PUBLISHING YOUR CONVERSATIONS:** Once you have completed your multi-turn interaction with a model and satisfied the final goal, you must generate a public link to share that exact conversation history.

- **HOW TO CREATE SHAREABLE LINK VIDEO**
- **Step 1:** Look for the **"Share"** or **"Publish Conversation"** button inside the specific AI model's interface (usually located at the top right or bottom of the chat window).
- **Step 2:** Click it to generate a public, shareable link. *Make sure the link privacy is set to public so our quality tools and reviewers can access it!*
- **Step 3:** Copy the link and paste it into the **"Model Published Conversation Link"** field on the main **Outlier** platform page before hitting final submission! Failing to paste the link will trigger our new blocking CoPilot linters.



5. Running the Conversation — Turns & Stop Condition

Tasks on this project are open-ended in length. You should carry the conversation for as many turns as necessary (**up to a maximum of 10 turns total**) to reach a fully satisfying answer that meets the **Final Goal to Reach**, then stop.

💡 **What if the model nails it on the very first response?** You cannot submit a single-turn task. Even if the first response is perfect, you must complete at least **1-3 required follow-up turns** to test how the model handles a continuation of the conversation.

The Three Golden Rules of Multi-Turn Chats

Rule 1: Document and Follow the Recommended Conversation Flow Your very first follow-up turn (Turn 2) must be based on the **Recommended Conversation Flow** list.

- **Core Flow Requirement:** You must use these recommended follow-up questions in order.
- **In the chat:** Copy the exact prompt in YOUR local language (for non-English contributors) from the list and paste it verbatim directly into your conversation with the models.
- **On the dashboard:** Copy and paste that exact same verbatim prompt in English into the field called "Selected next step."

Rule 2: Progressively Extend the Chat When Necessary After utilizing the mandatory initial sequence from the recommended flow, you can continue to guide the conversation if the project goal hasn't been fully satisfied.

- **Transition to Considerations:** If you have completely exhausted all of the prompts from the *Recommended Conversation Flow* list, you may then choose to use the remaining prompts taken from the **Considerations for Next Steps** list to continue.
- **Organic Guidance:** Alternatively, once the core flow is finished, you are free to write your own natural, custom prompts to guide the models. Your prompts must always adapt and respond to what the model actually said—engage with its real response instead of blindly re-pasting text.

Rule 3: Maintain Model Parity (Equal Patience) Treat both models exactly the same. Whatever patience, guidance, and number of turns you extend to Model A, you must also extend to Model B. Pushing one model through several follow-up turns while abandoning the other after a weak first turn creates an unfair comparison. How quickly or slowly each model reaches a satisfying answer is critical data—capture this behavior in your written feedback and let it inform your numerical scores.

6. Dimensions Scoring

For **each model**, score **all 11 dimensions** on a **1–3 scale**, judging the **whole conversation** (not any single turn).

Score Model A and Model B **independently**. Do not let one model's score influence the others - the comparison happens only at the final preference step.

The 11 dimensions fall into two groups: **Model Quality** (1–7) and **Product Experience** (8–11).

Model Quality Rubrics

Dimension	3 — No Issue	2 — Minor Issue	1 — Major Issue
1. Task Completion <i>Did the model help the user accomplish their goal? Whether the response meaningfully addresses the request, resolves the main intent, and moves toward completion.</i>	Fully satisfies the user's objective with little or no additional effort required.	Mostly satisfies the objective, but has small gaps, inefficiencies, or missing details that don't prevent the user from moving forward.	Only partially addresses the objective, misunderstands the core request, or leaves the user needing significant additional work.
2. Instruction Following <i>Did the model follow the user's explicit and implicit instructions? Includes format, tone, constraints, language, locale, length, style, and any requested do's or don'ts.</i>	Follows all explicit and implicit instructions, including format, tone, constraints, and subtle preferences.	Follows the main instructions, but misses one or more smaller constraints that don't significantly affect usefulness.	Ignores, contradicts, or fails to follow important instructions, making the output hard to use or misaligned with the request.
3. Accuracy and Groundedness <i>Is the response factually correct, reliable, and appropriately grounded? Covers factual accuracy, logical correctness, use of available context, and avoidance of unsupported claims.</i>	Accurate, reliable, well-supported, and appropriately caveated where needed.	Mostly accurate, but includes small unsupported claims, minor imprecision, or caveats that could be clearer.	Contains clear factual errors, hallucinations, misleading claims, or unsupported assertions that materially affect quality.

Dimension	3 — No Issue	2 — Minor Issue	1 — Major Issue
4. Relevance and Focus <i>Is the response focused on what the user actually asked?</i> Whether the model avoids unnecessary tangents, over-answering, boilerplate, or generic content.	Directly focused on the request; avoids unnecessary or irrelevant content.	Mostly relevant, but includes some extra, generic, or slightly unfocused content.	Largely off-topic, misses the actual intent, or focuses on the wrong parts of the request.
5. Clarity and Communication Quality <i>Is the response easy to understand and well structured?</i> Organization, readability, concision, wording, and whether it feels human-written rather than mechanical.	Clear, natural, well organized, concise, and easy to act on.	Understandable, but may be slightly wordy, awkward, poorly organized, or harder to scan than needed.	Confusing, poorly structured, overly verbose, unnatural, or difficult to understand or use.
6. Helpfulness and Practicality <i>Is the response genuinely useful in the user's real-world context?</i> Whether it gives actionable, practical, context-aware help rather than generic advice.	Highly useful, actionable, and tailored to the user's real-world context.	Useful overall, but could be more actionable, specific, or tailored to the situation.	Generic, impractical, hard to apply, or not meaningfully useful.
7. Locale and Cultural Fit <i>Does the response fit the user's language, locale, and cultural context?</i> Covers correct language/dialect, local terminology, cultural norms, units/currency/dates/form atting, region-specific expectations, and	Feels native, locally appropriate, and aligned with the user's language, region, and cultural context.	Understandable and mostly localized, but with small issues in phrasing, terminology, formatting, or cultural fit.	Uses the wrong language or dialect, feels unnatural for the locale, applies incorrect regional assumptions, or is culturally inappropriate.

Dimension	3 — No Issue	2 — Minor Issue	1 — Major Issue
avoidance of unnatural translations. Especially important for multilingual and region-specific evals.			

Product Experience Rubrics

Dimension	3 — No Issue	2 — Minor Issue	1 — Major Issue
8. Conversation Flow <i>Does the assistant maintain a smooth, useful conversation over multiple turns?</i> Whether it remembers context, builds on prior turns, avoids making the user repeat themselves, and advances naturally.	Maintains context smoothly, builds on prior turns, and moves the interaction forward naturally.	Mostly maintains flow, but has small moments of repetition, missed nuance, or slightly awkward turn-taking.	Loses track of the conversation, forces unnecessary repetition, or fails to advance the interaction usefully.
9. Friction and Effort Required <i>How much extra work does the user need to do after receiving the response?</i> Whether the experience feels efficient or frustrating.	The user can act immediately with little or no additional clarification, editing, or effort.	The user needs small clarifications, edits, or adjustments before the response is usable.	The user needs significant additional work, must re-explain the request, heavily correct the output, or restart the task.
10. Multi-Turn Consistency <i>Does the assistant consistently maintain and apply the user's context, preferences, and intent across the conversation?</i> Rewards using relevant info from earlier turns, continuity in tone and direction, and not making the user repeat themselves. Penalizes	Consistently carries forward context, preferences, constraints, and goals across turns in a helpful, natural way.	Maintains the main context and intent, but misses some smaller prior details or shows minor gaps in continuity.	Forgets or contradicts important prior context, ignores stated preferences/constraints, makes inappropriate assumptions, or requires the user to repeat key information.

Dimension	3 — No Issue	2 — Minor Issue	1 — Major Issue
contradicting prior context, forgetting stated preferences, or over-assuming.			
11. Delight and Overall Experience <i>Does the interaction feel satisfying, polished, and pleasant?</i> The softer product-quality layer: tone, empathy, naturalness, and whether the assistant feels like something users would want to keep using.	Polished, natural, satisfying, and pleasant; appropriate tone; easy to continue using.	Usable and generally pleasant, but may feel flat, robotic, impersonal, or slightly rough around the edges.	Frustrating, awkward, unpleasant, or meaningfully below the quality users expect from a helpful assistant.

7. How to Write Strong Feedback

You write **two** feedback blocks (one per model, independent) plus **one** final justification (comparison). **Ensure that in the feedback you write you are talking about how the model behaved in response to the next steps you gave it alongside if the final goal was reached.**

Element	Guidance
Length	4 - 6 sentences. Long enough to dive into what worked and what didn't, short enough to stay focused.
Specificity	Cite concrete examples from the conversation. "The third turn contradicted the constraint set in turn one" beats "consistency was off."
Anchor to the criteria	Tie your points to the dimensions above - your written feedback and your 1–3 scores should tell the same story. If you scored a dimension a 1, the feedback should explain why.
Consideration of Next Steps	Talk about how the model behaved and what worked and what didn't work when you

Element	Guidance
	attempted to guide it using one of the next steps provided. Reference specific examples from the model response.
Turn behavior	Note how many turns it took and how the model handled follow-ups - how quickly it reached the final goal is real signal.
Independence (Model A & B feedback)	Judge each response on its own merits. Do not write "Model A was better than Model B" in the per-model feedback - save all comparison for the final justification.
Comparison (final justification only)	The final justification is the only place you compare. Name the concrete differences that drove your preference, and acknowledge where the model you didn't pick was actually stronger.
FINAL GOAL	When writing your justifications, you MUST address whether the model reached the final goal, how fully it got there, and what worked or fell short along the way.

Example:

Prompt:

Can you write an artist bio about Kim Myhre Clarke using the information I provide?

Considerations for Next Steps:

1. The user might ask the AI to elaborate on a specific aspect of Kim's artistic process, such as her use of glazes or engobes.
2. The user might request the AI to highlight Kim's world travels and how they influenced her work.
3. The user might ask the AI to focus more on the influence of specific artists like Frida Kahlo or Tim Burton on Kim's style.

4. The user might ask the AI to write a shorter or longer version of the bio for different purposes.
5. The user might ask the AI to suggest ways to incorporate images of Kim's work into the bio.

Final Goal to Reach:

The user wants the AI to write a compelling artist bio for Kim Myhre Clarke, incorporating details about her life, artistic background, influences, creative process, and recent work, based on the provided text.

Example Strong per-model feedback (independent)

The model reached the final goal by the second turn, producing a compelling bio that pulled Kim's background, influences, and creative process straight from the text I provided rather than inventing details (Task Completion, Accuracy and Groundedness). For my first follow-up I used one of the next-step considerations and asked it to highlight her world travels; it wove that in cleanly without dropping the earlier detail about her glaze and engobe work (Multi-Turn Consistency). The structure was clean and the tone read like a professional artist bio, so it was usable with little editing (Clarity, Instruction Following). The one weakness was a generic closing paragraph that added length without new information, slightly diluting the focus (Relevance and Focus).

Why this works: it opens by stating whether and how well the final goal was met, cites specific detail from the actual conversation (the travel follow-up, the glaze work, the weak closing), anchors each point to a named dimension, notes the turn count and the required Considerations follow-up, and stays fully independent - no mention of the other model. Four to five tight sentences, not a paragraph.

Example Strong final justification (comparison)

Model A reached a fully satisfying bio in two turns, while Model B needed four and still left the creative-process section thin, so Model A met the final goal more completely. Model A also kept the bio firmly grounded in the text I provided, whereas Model B invented an exhibition that wasn't in my notes, which undercut its accuracy. Model B did handle the world-travel follow-up with richer, more vivid detail, which was a genuine strength. But the user's goal was a compelling bio built faithfully from the provided information, and Model A delivered that more reliably and with less back-and-forth. For its stronger groundedness and faster goal completion, I preferred Model A.

Why this works: comparison lives only here (never in the per-model feedback), it names the concrete differences that drove the choice, ties the decision back to the user's actual final goal, and honestly credits where the model that wasn't chosen was stronger.

Weak feedback (avoid)

Pretty good response. Helpful and well-written. Better than the other one.

Why it fails: vague with no specifics from the conversation, never states whether or how well the final goal was reached, not anchored to any dimension, and comparative where it should be independent.

8. Common Mistakes to Avoid

- **Vague feedback.** "Good response" tells us nothing. Cite specifics from the conversation.
- **Comparing in the wrong place.** Model A and Model B feedback must be independent. Comparison belongs only in the final justification.
- **Unequal patience.** Running one model through several follow-ups while cutting the other short breaks parity. Treat both the same.
- **Skipping the required first follow-up.** Your first follow-up must come from the Recommended Conversation Flow and Considerations for Next Steps.
- **Stopping before / overshooting the goal.** Use the **final goal** as your stop condition - don't quit early, and don't pad with extra turns once it's satisfied.
- **Scores that don't match the feedback.** If you scored a dimension low, your written feedback should say why. They must agree.
- **Writing feedback in the wrong language.** The prompt is in your language; **all feedback and justifications are in English.**
- **Editing the preset prompt.** Paste it as delivered.